

STA 630 - Homework 4

Anthony Bernardi

April 9th, 2025

1 Problem 1

Consider the balanced, additive, one-way ANOVA model:

$$Y_{ij} = \mu + \alpha_i + \epsilon_{ij} \quad (1)$$

$$i = 1, \dots, I, j = 1, \dots, J \quad (2)$$

Where epsilon is normally distributed with parameters $0, \sigma_\epsilon^2$ iid, with μ, α_i in the real numbers, and σ_{ϵ}^2 a non-negative number. We adopt the prior structure that is the product of independent conjugate priors, where μ is a flat prior, alpha is normally distributed, and sigma follows the inverse-gamma distribution with parameters a and b, as we have seen.

We will assume that σ_α^2 is known, representing the variance of the alpha parameter. We assume a and b for the Inverse Gamma are known as well.

1.1 Part a

Derive the full condition distributions for μ, α_i , and σ_ϵ^2 , as required to implement the Gibbs sampler.

We can copy this from the hand-written sheet later...

1.2 Part b

What is meant by convergence diagnosis? Describe some tools you might use to assist in this regard, what might you do to improve a sampler that is converging slowly?

I would think that, given a Gibbs sampler that is converging slowly, both the starting point of the Gibbs sampling algorithm along with the autocorrelation of the samples would be two factors I would look at first. Given that this is a class of MCMC methods, and these are not independent samples, your starting point matters immensely in evaluating and running your algorithm. Each sample obviously depends on the prior for the full conditionals, so I would evaluate the initial value and starting point when diagnosing a specific Gibbs sampler.

Looking into the autocorrelation or "stickiness" to use Hoff's term, is what I would evaluate next. If the sampler oscillates between a few values, and doesn't progress throughout the posterior density curve or "mix" well, a sampler with high autocorrelation may be causing this.

As far as fixing this, picking a different starting point that causes the sampler to mix better may prove beneficial.

1.3 Part c

Suppose that σ_ϵ^2 is known. What conditions on the data or the prior might lead to slow convergence for the other parameters via the Gibbs sampler?

I imagine that, if the variance is known, and we are estimating one fewer parameter via Gibbs sampling, we might be "at the mercy" of the now-known variance value, so to speak. If the variance is very small, and the data is very close to the mean, we might have a situation where the posterior density is very peaked around the mean. This would lead to slow convergence, as the sampler would be "stuck" in a small region of the parameter space.

Largely speaking, a known variance value could be unhelpful in estimating the posterior, in that this might result in different distribution shapes, which could lead to the sampling algorithm being stuck, ultimately obviously slowing convergence.

1.4 Part d

Some scratch work before completing the Gibbs sampler can also be found in the hand-written notes.

The R code written in implementing the Gibbs sampler is as follows.

```
# comparing a gibbs sampler for the centered and un-centered algos
## loading the data
```

```

y <- rnorm(25, 1, 1)
mean.y <- mean(y)
var.y <- var(y)
n <- length(y)

# above serve as likelihoods. Now to the gibbs sampling
# starting values
S <- 1000
phi <- matrix(nrow=S, ncol=2)
phi[1,] <- c(mean.y, 1/var.y)
lambda <- matrix(nrow=S, ncol=2) # for the centered case

# Gibbs sampling algo
set.seed(1870)
for (s in 2:S){
  # generating the new mean value
  i <- 1
  mu_n <- rnorm(1, phi[i,1], phi[i,2] / S) # in the matrix
  i <- i + 1
  phi[i+1,1] <- mu_n

  # other parameter
  i <- 1
  alpha_n <- rnorm(1, phi[i+1,1], phi[i,2]) # other parameter
  phi[i+1,2] <- alpha_n
}

# we can now run this and proceed with the algorithm,
# now we will handle the centered case
# working with our new parameter matrix
lambda <- matrix(nrow=S, ncol=2) # for the centered case

lambda[1,] <- c(mean.y, mean.y + 1/var.y)

for (s in 2:S){
  # generating the new mean as before
  i <- 1
  mu_n <- rnorm(1, lambda[i,1], lambda[i,2] / S)

```

```

i <- i + 1
lambda[i+1, 1] <- mu_n

# other parameter
i <- 1
eta_n <- rnorm(1, lambda[i+1, 1], lambda[i,2])
lambda[i+1,2] <- eta_n
}

```

As for which performs better, the parameter case where there is centering compared with where there is not, I would think that the un-centered case performs better, as the centered case is prone to stickiness and limiting its ability to approximate the posterior.

2 Problem 2